Missing Values and Plotting Pretty Graphs

by Alex Dance

Purpose

Share how cleaned data and how present data

Agenda

- Cleaning
- Presenting
- Comparing

- Not Work
- Not Work Well
- Not Work Well
- Work Just OK
- Work Well

Not Showed as we have all done in exercises

- Profile report
- Standard Correlations

Stats Lab

1.1.4 Lab 1.1.4 Statistics - part 2 l

```
Rounded Version is:  y = 0.04 x +  31.55

<matplotlib.collections.PathCollection at 0x1daaec117c8>
```

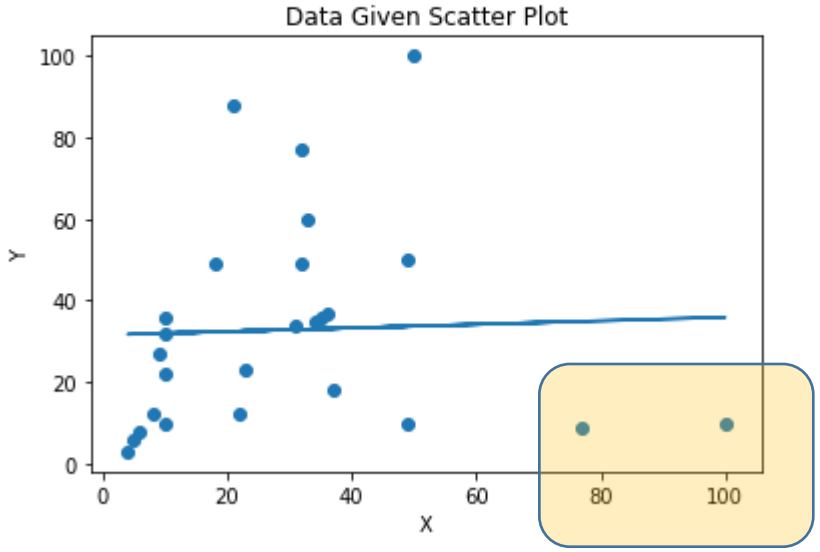


```
Rounded Version is:  y = 0.25 x +  28.84

Out[14]: <matplotlib.collections.PathCollection at 0x26b2db355c8>
```

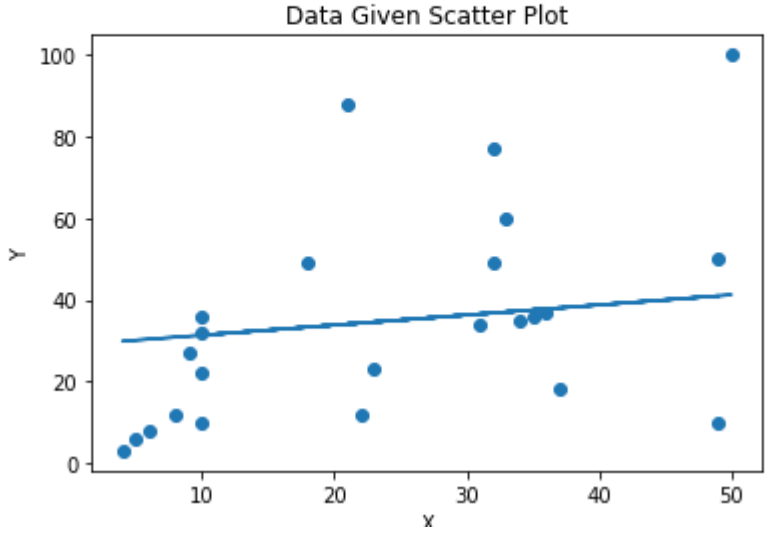

Presenting

Wanted to print some impressive graphs

https://seaborn.pydata.org/generated/seaborn.kdeplot.html

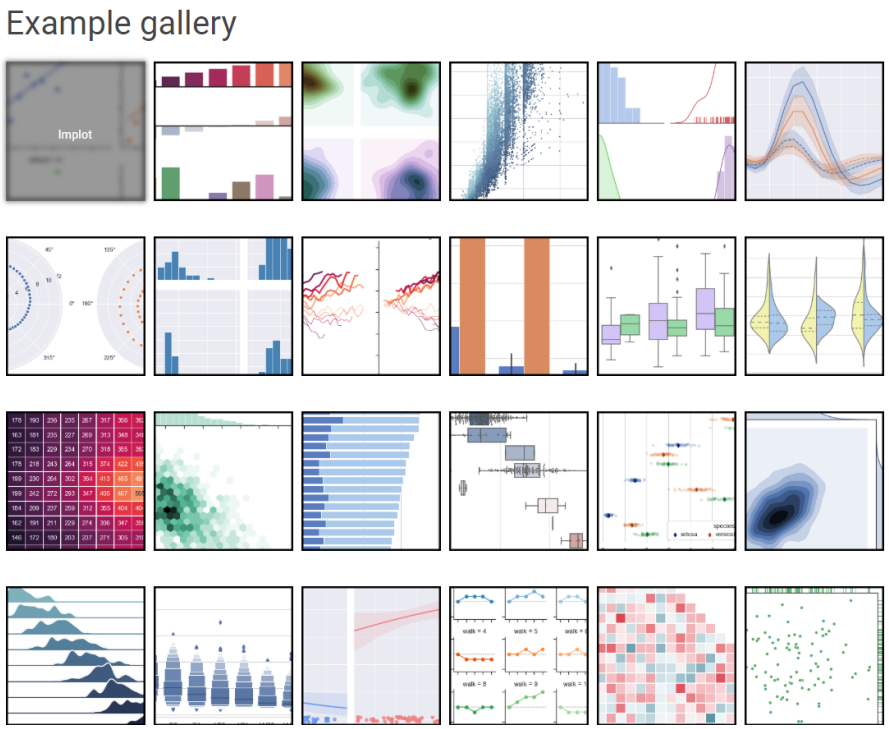

Dataset of Melbourne Data and Model Comparison from Kaggle

Started with some basic graphs

Price of 2 beds

```
In [68]: ax = sns.kdeplot(adjustbedtwo.Price)
```

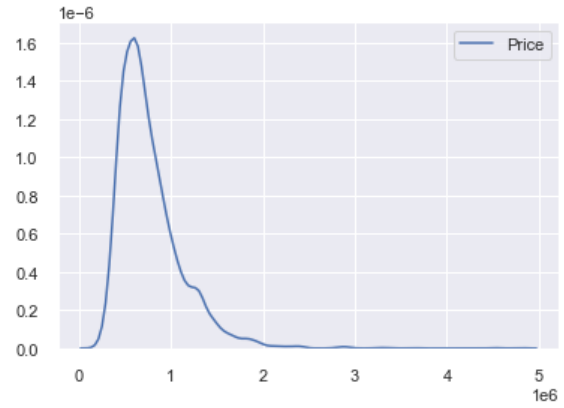

Log of the same thing

```
ax = sns.kdeplot(np.log(adjustbedtwo.Price))
```

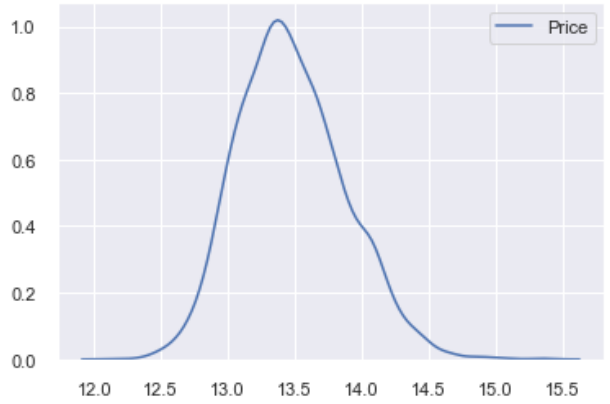

Distance of 2 beds

```
In [71]: ax = sns.kdeplot(adjustbedtwo.Distance)
```

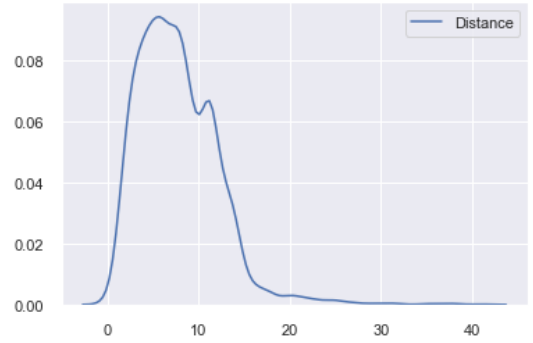


```
In [72]: plt.scatter(adjustbedtwo['Price'], adjustbedtwo['Distance'])
plt.show()
```

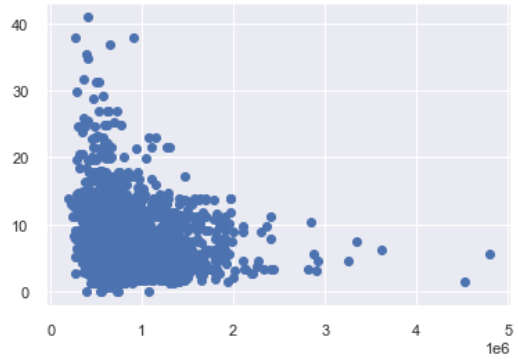

Wanted a good KDE Plot

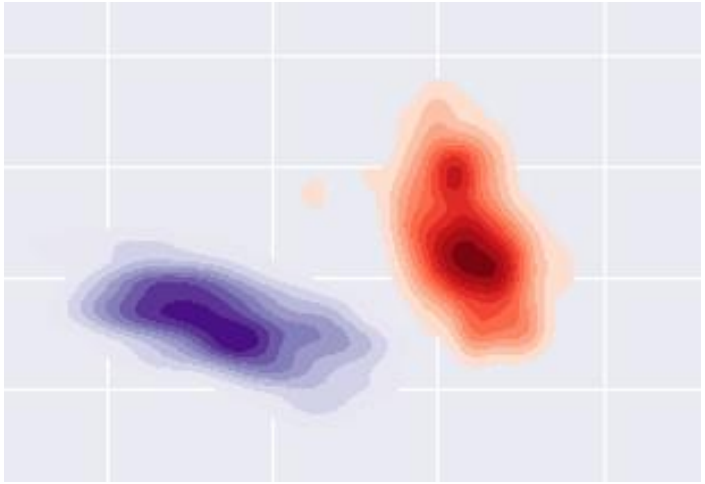

Wanted

```
108]: # Distance
      ax = sns.kdeplot(beedone.Price, beedone.Distance,cmap="Reds", shade=True, shade_lowest=False,clip=(100,500000))
      ax = sns.kdeplot(beedtwo.Price, beedtwo.Distance,cmap="Blues", shade=True, shade_lowest=False, clip=(100,500000))
```

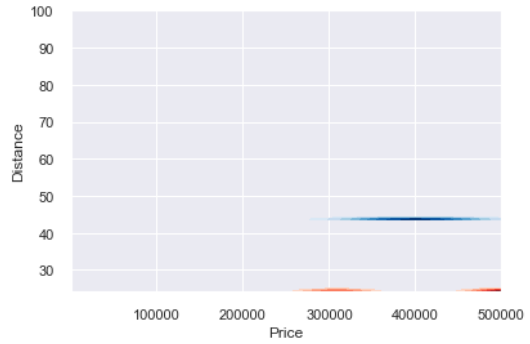

Got

Started
Fixing
the Data

Dropping

- reduced_X_train = X_train.drop(cols_with_missing, axis=1)
- https://www.w3resource.com/pandas/dataframe/dataframe-drop.php

DataFrame - drop() function

The drop() function is used to drop specified labels from rows or columns.

Remove rows or columns by specifying label names and corresponding axis, or by specifying directly index or column names. When using a multi-index, labels on different levels can be removed by specifying the level.

Syntax:

```
]: print(cols_with_missing)

['Car', 'BuildingArea', 'YearBuilt']
```

```
]: # Drop columns in training and validation data
reduced_X_train = X_train.drop(cols_with_missing, axis=1)
reduced_X_valid = X_valid.drop(cols_with_missing, axis=1)
```

Excluding

alexmelbexclude= alexmelb.select_dtypes(exclude=['object'])

```
In [135]: alexmelb.head()
```

Out[135]:

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	Distance	Postcode	...	Car	Landsize	BuildingArea
0	Abbotsford	85 Turner St	2	h	1480000.0	S	Biggin	3/12/2016	2.5	3067.0	...	1.0	202.0	NaN
1	Abbotsford	25 Bloomburg St	2	h	1035000.0	S	Biggin	4/02/2016	2.5	3067.0	...	0.0	156.0	79.0
2	Abbotsford	5 Charles St	3	h	1465000.0	SP	Biggin	4/03/2017	2.5	3067.0	...	0.0	134.0	150.0
3	Abbotsford	40 Federation La	3	h	850000.0	PI	Biggin	4/03/2017	2.5	3067.0	...	1.0	94.0	NaN
4	Abbotsford	55a Park St	4	h	1600000.0	VB	Nelson	4/06/2016	2.5	3067.0	...	2.0	120.0	142.0

5 rows × 22 columns

```
In [25]: alexmelbexclude= alexmelb.select_dtypes(exclude=['object'])
```

```
In [26]: alexmelbexclude.head()
```

Out[26]:

	Rooms	Price	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	BuildingArea	YearBuilt	Lattitude	Longtitude	Proper
0	2	1480000.0	2.5	3067.0	2.0	1.0	1.0	202.0	NaN	NaN	-37.7996	144.9984	
1	2	1035000.0	2.5	3067.0	2.0	1.0	0.0	156.0	79.0	1900.0	-37.8079	144.9934	
2	3	1465000.0	2.5	3067.0	3.0	2.0	0.0	134.0	150.0	1900.0	-37.8093	144.9944	

The Simple Imputer Library is an easy option

https://scikit-learn.org/stable/modules/generated/sklearn.impute.SimpleImputer.html

```
>>> import numpy as np
>>> from sklearn.impute import SimpleImputer
>>> imp_mean = SimpleImputer(missing_values=np.nan, strategy='mean')
>>> imp_mean.fit([[7, 2, 3], [4, np.nan, 6], [10, 5, 9]])
SimpleImputer()
>>> X = [[np.nan, 2, 3], [4, np.nan, 6], [10, np.nan, 9]]
>>> print(imp_mean.transform(X))
[[ 7.   2.   3. ]
 [ 4.   3.5  6. ]
 [10.   3.5  9. ]]
```

Methods

`fit`(self, X[, y])	Fit the imputer on X.
`fit_transform`(self, X[, y])	Fit to data, then transform it.
`get_params`(self[, deep])	Get parameters for this estimator.
`set_params`(self, **params)	Set the parameters of this estimator.
`transform`(self, X)	Impute all missing values in X.

I used it myself

```
In [29]: alexmelbadjustnohead = pd.DataFrame(my_imputer.fit_transform(alexmelbexclude))
```

```
In [30]: alexmelbadjustnohead.head()
```

Out[30]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	2.0	1480000.0	2.5	3067.0	2.0	1.0	1.0	202.0	151.96765	1964.684217	-37.7996	144.9984	4019.0
1	2.0	1035000.0	2.5	3067.0	2.0	1.0	0.0	156.0	79.00000	1900.000000	-37.8079	144.9934	4019.0
2	3.0	1465000.0	2.5	3067.0	3.0	2.0	0.0	134.0	150.00000	1900.000000	-37.8093	144.9944	4019.0
3	3.0	850000.0	2.5	3067.0	3.0	2.0	1.0	94.0	151.96765	1964.684217	-37.7969	144.9969	4019.0
4	4.0	1600000.0	2.5	3067.0	3.0	1.0	2.0	120.0	142.00000	2014.000000	-37.8072	144.9941	4019.0

Initially it dropped Titles

```
In [31]: alexmelbadjust = pd.DataFrame(my_imputer.fit_transform(alexmelbexclude),columns = alexmelbexclude.columns)
```

```
In [32]: alexmelbadjust.head()
```

Out[32]:

	Rooms	Price	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	BuildingArea	YearBuilt	Lattitude	Longtitude	Propertycount
0	2.0	1480000.0	2.5	3067.0	2.0	1.0	1.0	202.0	151.96765	1964.684217	-37.7996	144.9984	4019.0
1	2.0	1035000.0	2.5	3067.0	2.0	1.0	0.0	156.0	79.00000	1900.000000	-37.8079	144.9934	4019.0
2	3.0	1465000.0	2.5	3067.0	3.0	2.0	0.0	134.0	150.00000	1900.000000	-37.8093	144.9944	4019.0
3	3.0	850000.0	2.5	3067.0	3.0	2.0	1.0	94.0	151.96765	1964.684217	-37.7969	144.9969	4019.0
4	4.0	1600000.0	2.5	3067.0	3.0	1.0	2.0	120.0	142.00000	2014.000000	-37.8072	144.9941	4019.0

Then Presenting Got Better

Started Getting Somewhere

```
In [139]: ax = sns.kdeplot(np.log(adjustbedone.Price), np.log(adjustbedone.Distance),cbar=True)
```

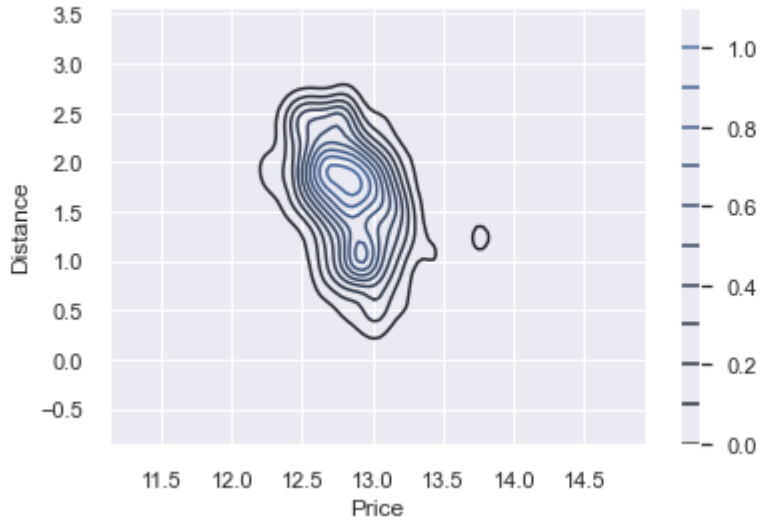

Not Work

In [83]:
```
ax = sns.kdeplot(adjustbedone.Price, adjustbedone.Distance,cmap="Reds", shade=True, shade_lowest=False)
ax = sns.kdeplot(adjustbedtwo.Price, adjustbedtwo.Distance,cmap="Greys", shade=True, shade_lowest=False)
ax = sns.kdeplot(adjustbedthree.Price, adjustbedthree.Distance,cmap="Blues", shade=True, shade_lowest=False)
ax = sns.kdeplot(adjustbedfour.Price, adjustbedfour.Distance,cmap="Greens", shade=True, shade_lowest=False)
ax = sns.kdeplot(adjustbedfive.Price, adjustbedfive.Distance,cmap="Purples", shade=True, shade_lowest=False)
```

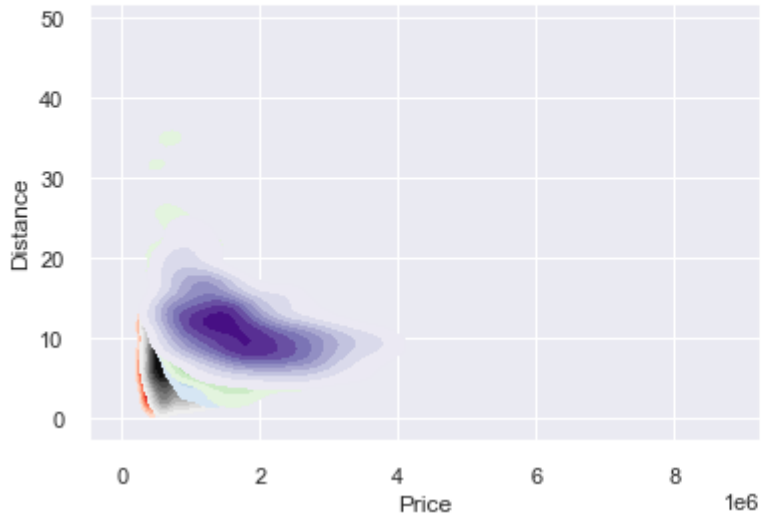

Graphs started Looking OK

```
In [77]: ax = sns.kdeplot(adjustbedtwo.Price, adjustbedtwo.Distance,
```

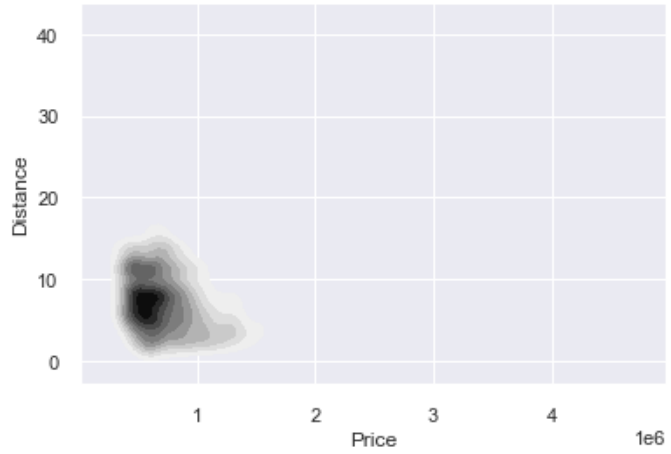


```
In [90]: ax = sns.kdeplot(np.log(adjustbedtwo.Price), adjustbedtwo.Distance,cm
```

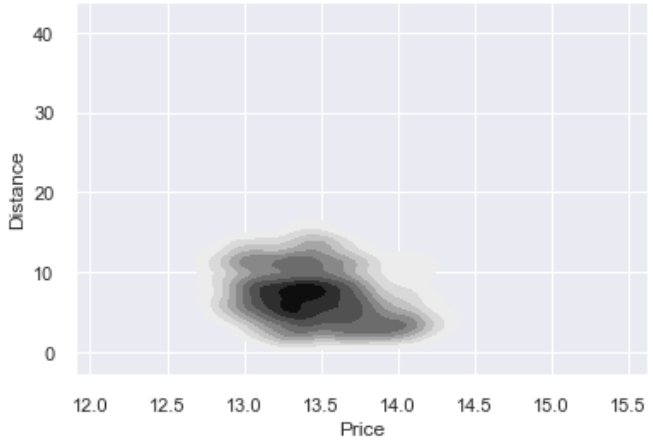

I like looking at them

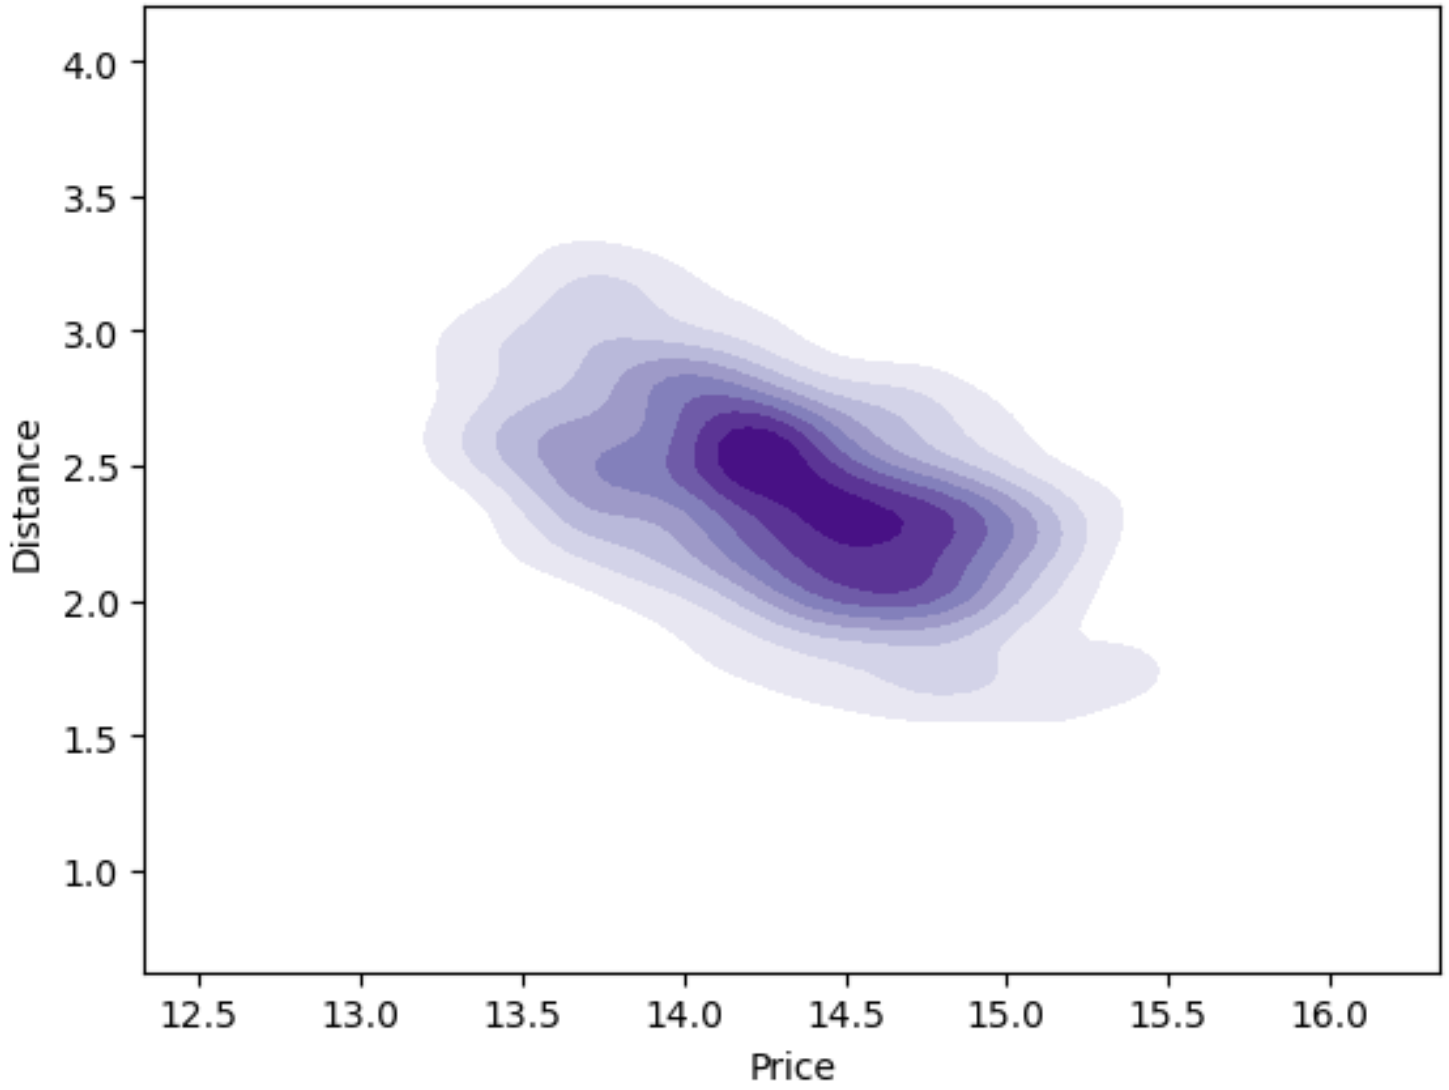

Works best with limited options – 1 bed Vs 5 Bed

1 Bed Vs 5 Bed

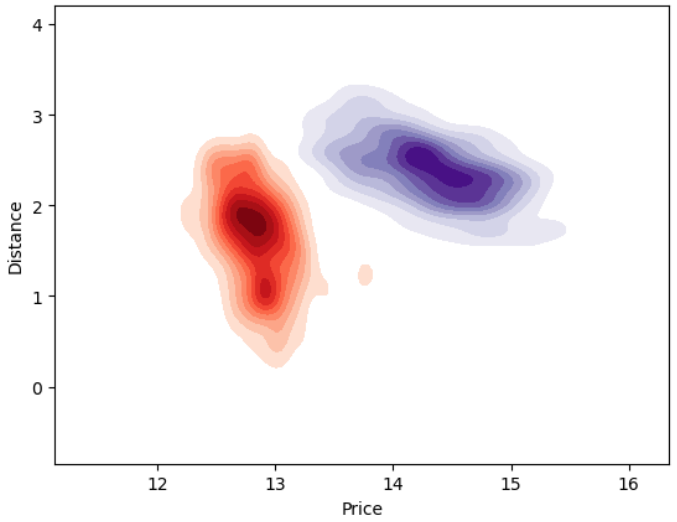

1 Bed Vs 3 bed Vs 5 Bed

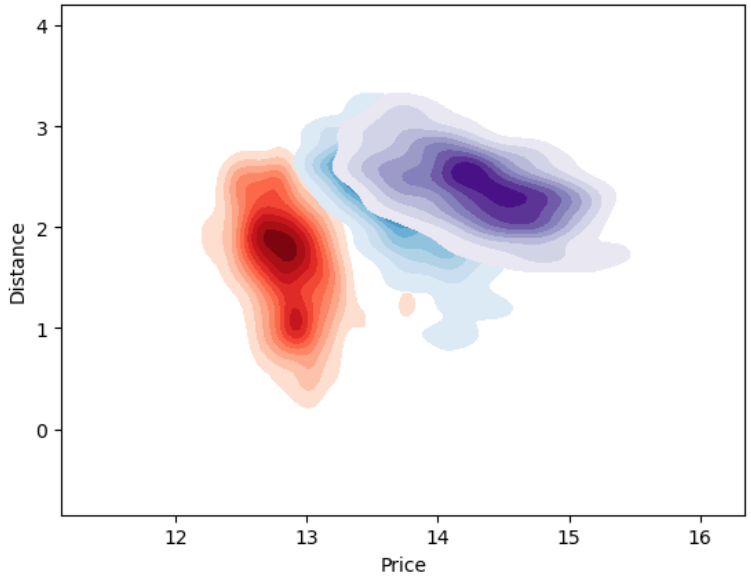

Graphs Comparison

```
ax = sns.kdeplot(np.log(adjustbedone.Price), np.log(adjustbedone.Distance),cma
```

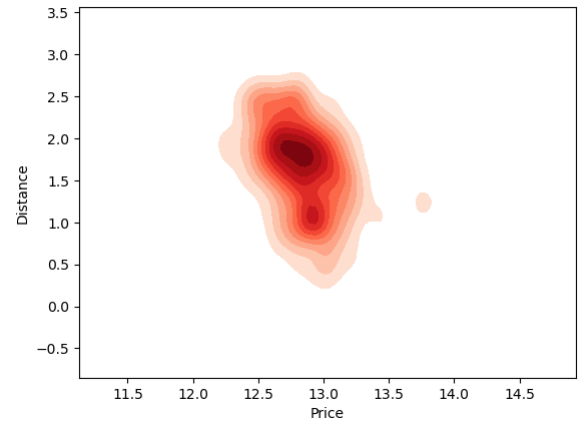


```
ax = sns.kdeplot(np.log(adjustbedtwo.Price), adjustbedtwo.Distance,cmap="Gr
```

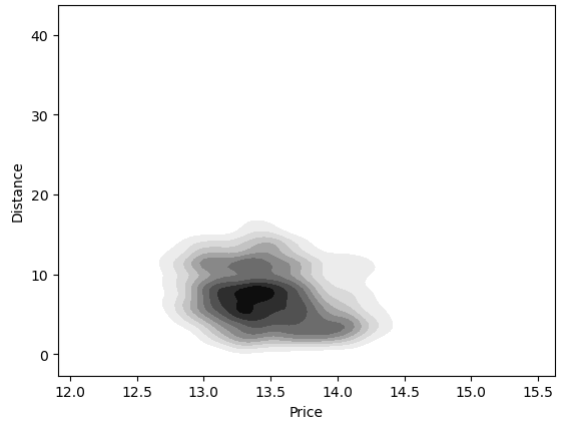


```
ax = sns.kdeplot(np.log(adjustbedthree.Price), np.log(adjustbedthree.Distanc
```

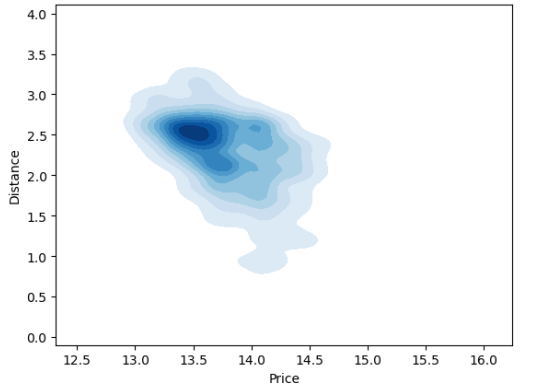


```
ax = sns.kdeplot(np.log(adjustbedfour.Price), np.log(adjustbedfour.Distance
```

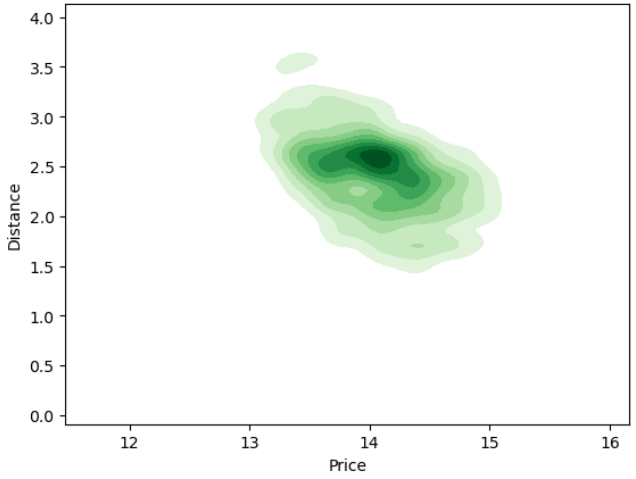

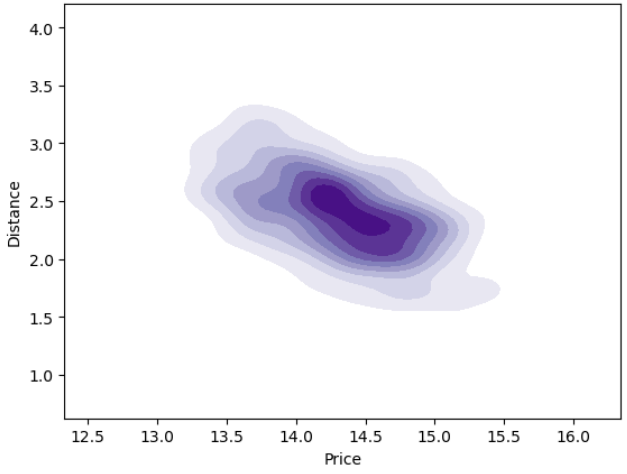

Useful

Found how to split data

```
# Divide data into training and validation subsets
X_train, X_valid, y_train, y_valid = train_test_split(X, y, train_size=0.8, test_size=0.2,random_state=0)
```

Useful resource

https://www.machinelearningplus.com/plots/top-50-matplotlib-visualizations-the-master-plots-python/

22. Density Plot

Density plots are a commonly used tool visualise the distribution of a continuous variable. By grouping them by the 'response' variable, you can inspect the relationshi between the X and the Y. The below case if for representational purpose to describe how the distribution of city mileage varies with respect the number of cylinders.

Show Codes

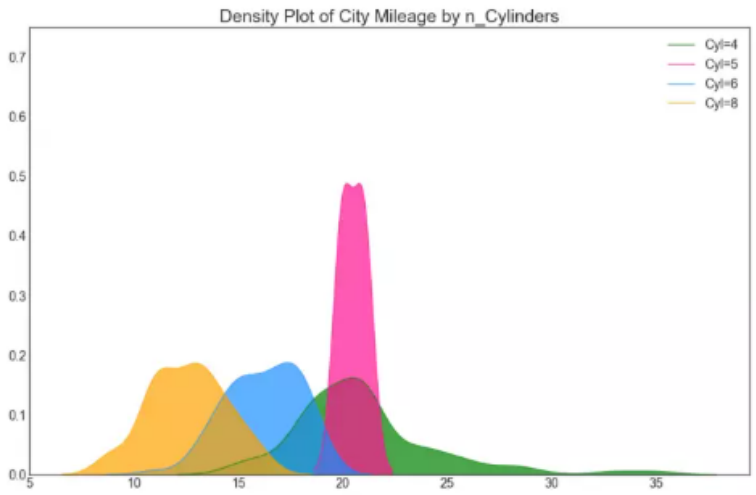

Show Code

```
# Import Data
df = pd.read_csv("https://github.com/selva86/datasets/raw/master/mpg_ggplot2.csv")

# Draw Plot
plt.figure(figsize=(16,10), dpi= 80)
sns.kdeplot(df.loc[df['cyl'] == 4, "cty"], shade=True, color="g", label="Cyl=4", alpha=.7)
sns.kdeplot(df.loc[df['cyl'] == 5, "cty"], shade=True, color="deeppink", label="Cyl=5", alpha=.7)
sns.kdeplot(df.loc[df['cyl'] == 6, "cty"], shade=True, color="dodgerblue", label="Cyl=6", alpha=.
sns.kdeplot(df.loc[df['cyl'] == 8, "cty"], shade=True, color="orange", label="Cyl=8", alpha=.7)

# Decoration
plt.title('Density Plot of City Mileage by n_Cylinders', fontsize=22)
plt.legend()
plt.show()
```

Worked well for Melbourne House Price Data

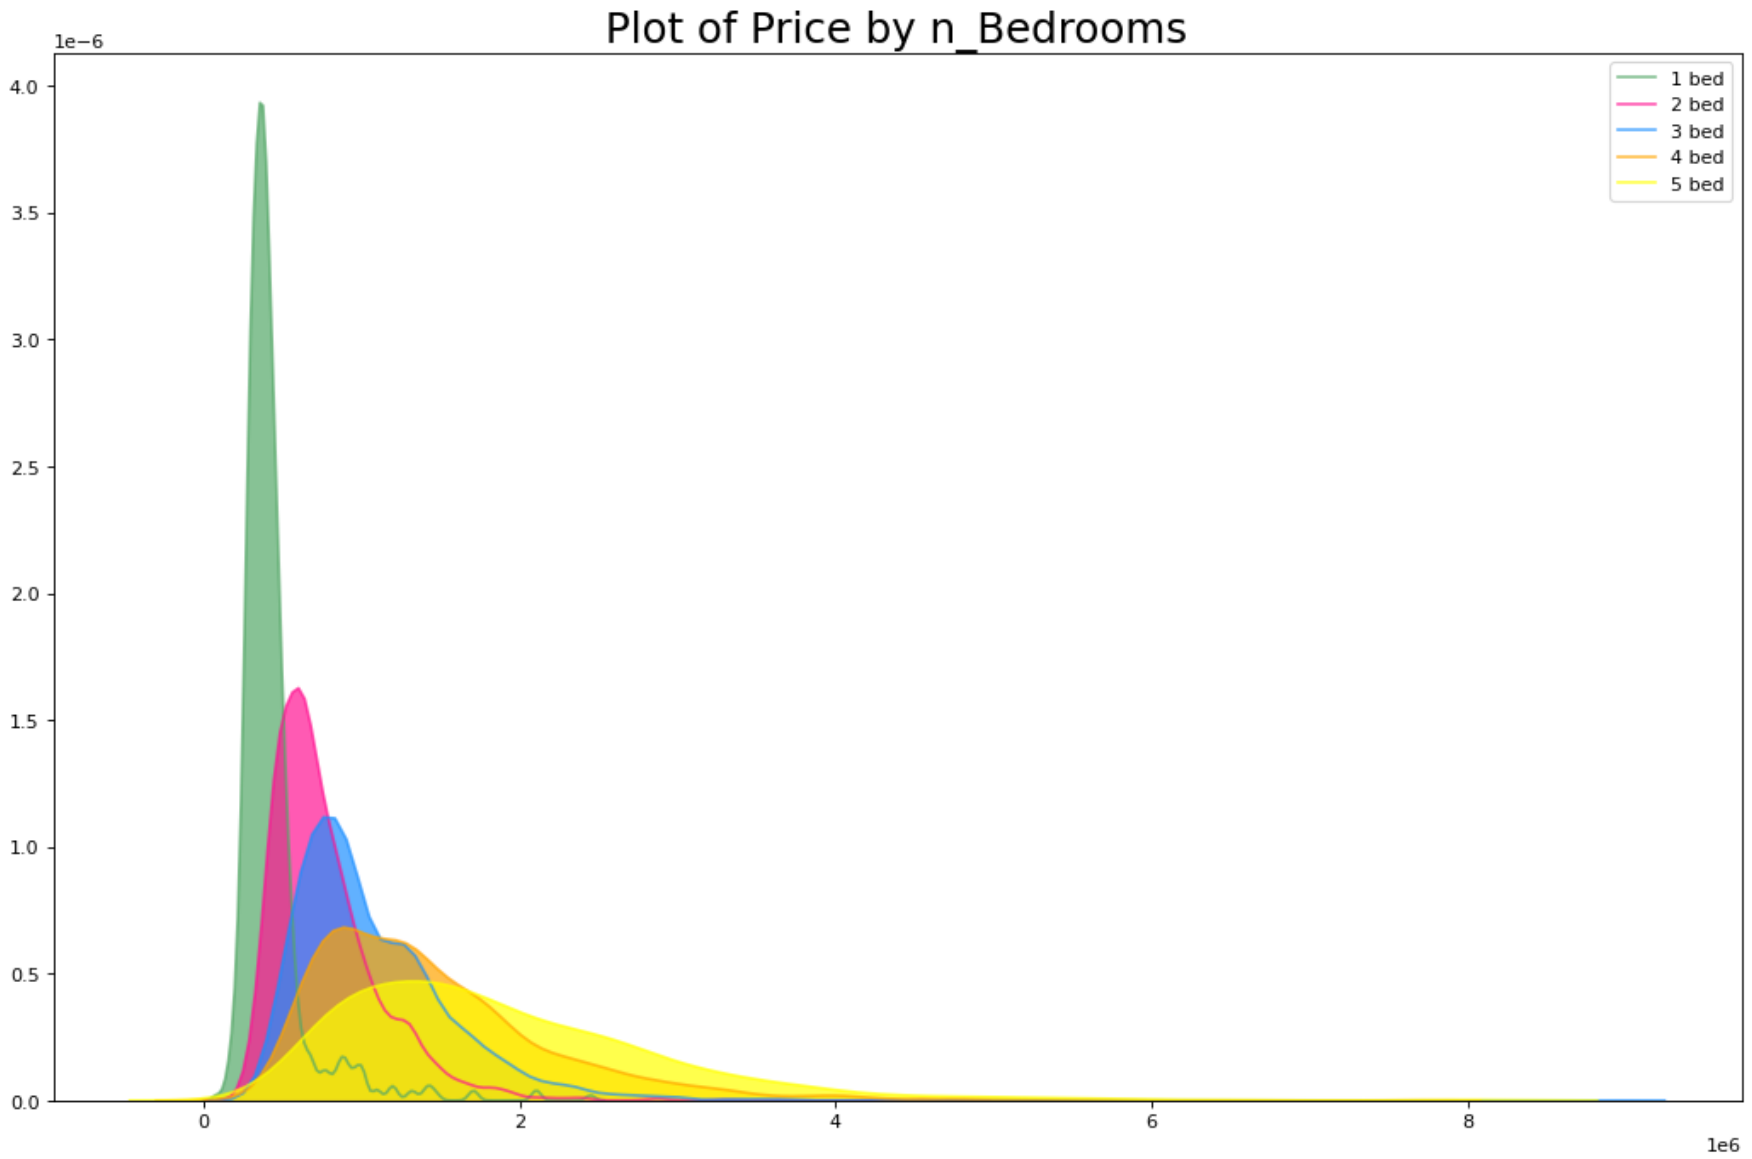

Useful resource

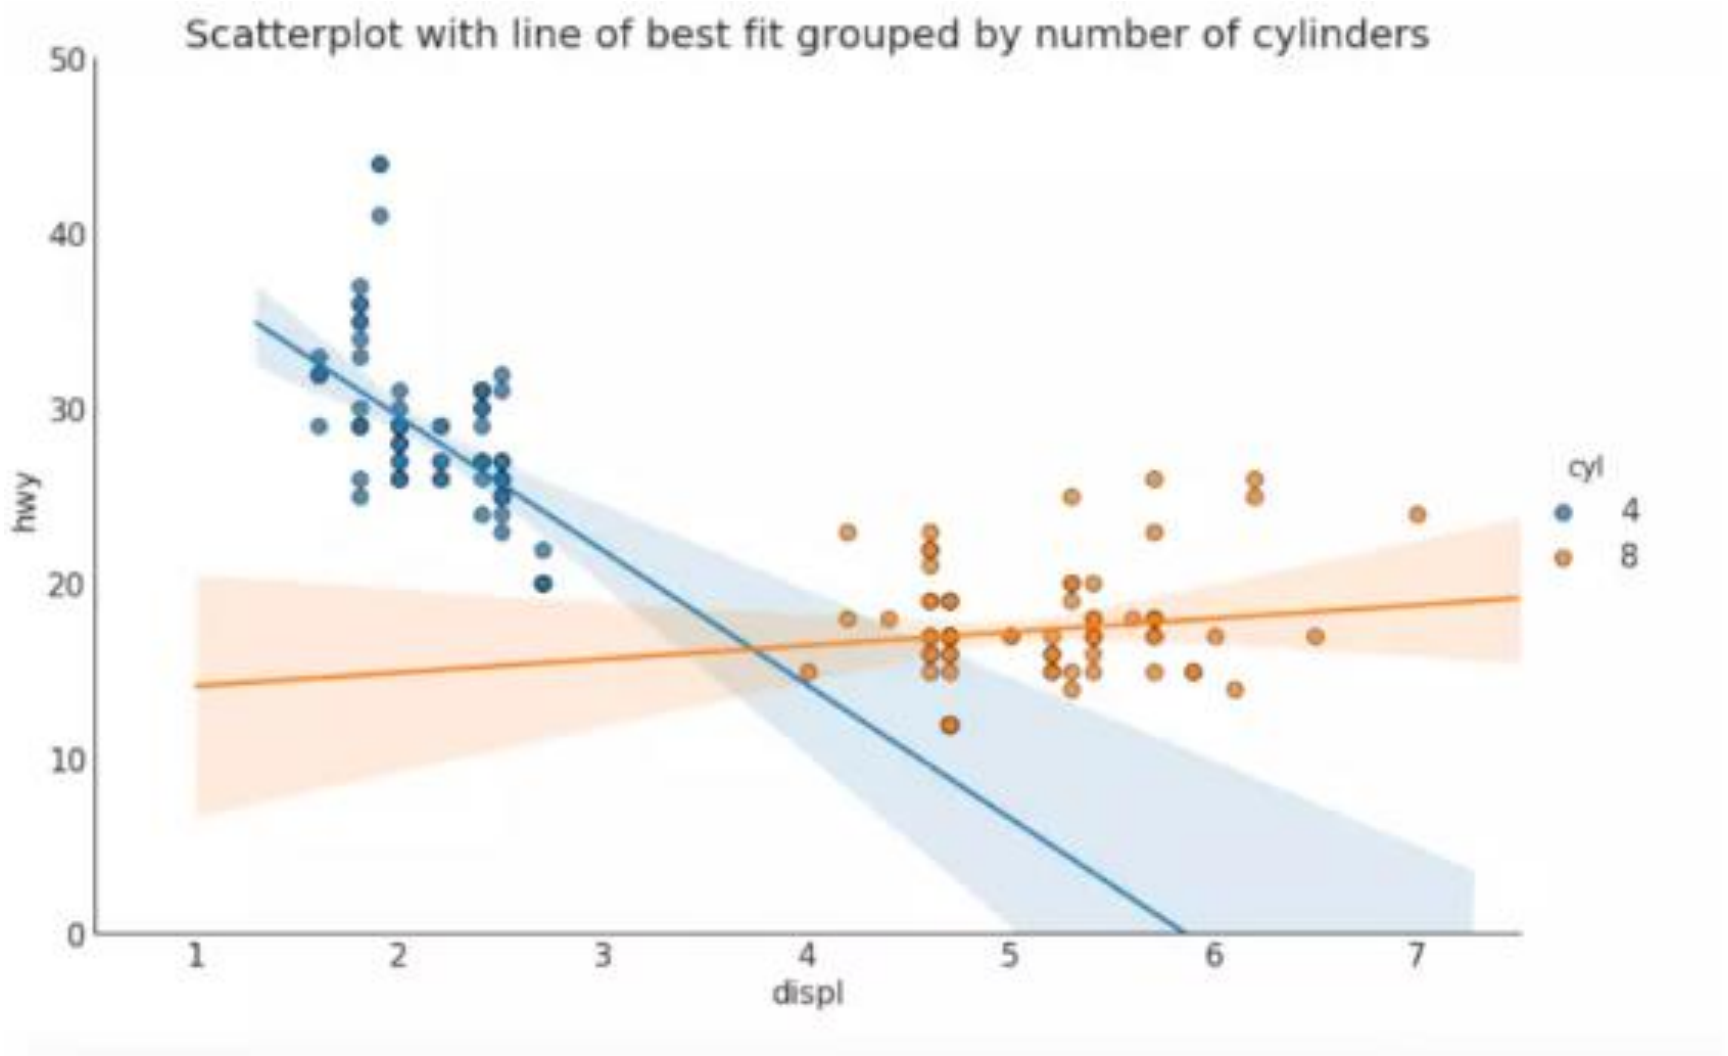

Useful resource

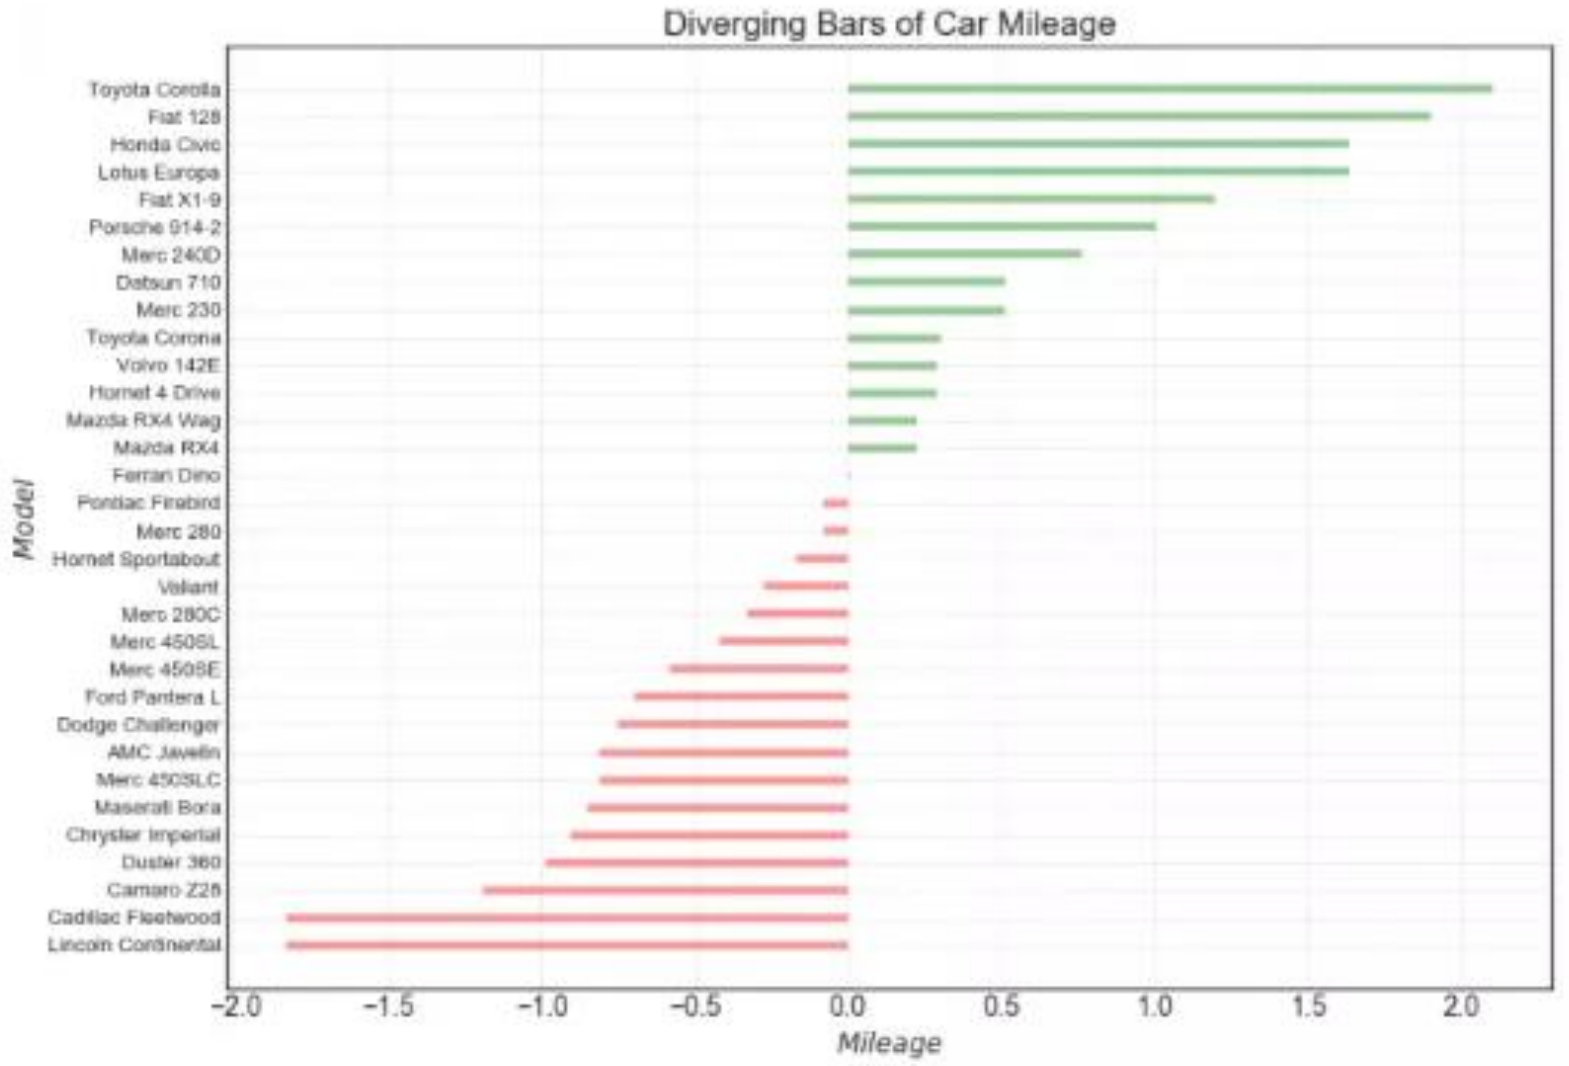

Useful resource

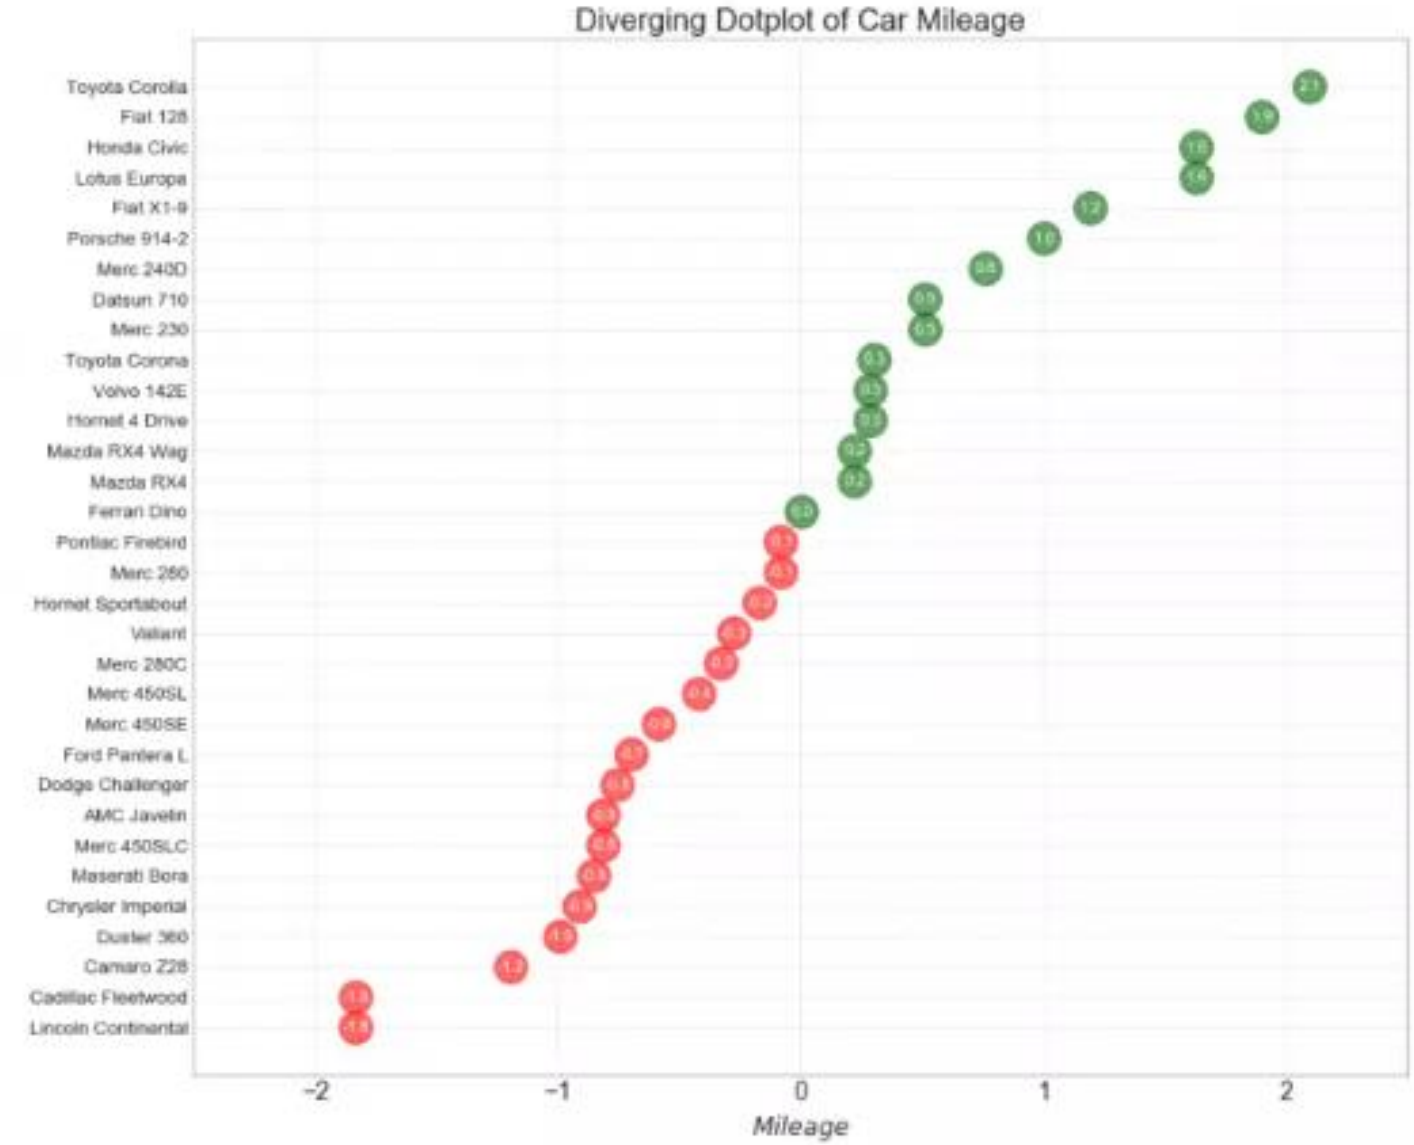

Useful resource

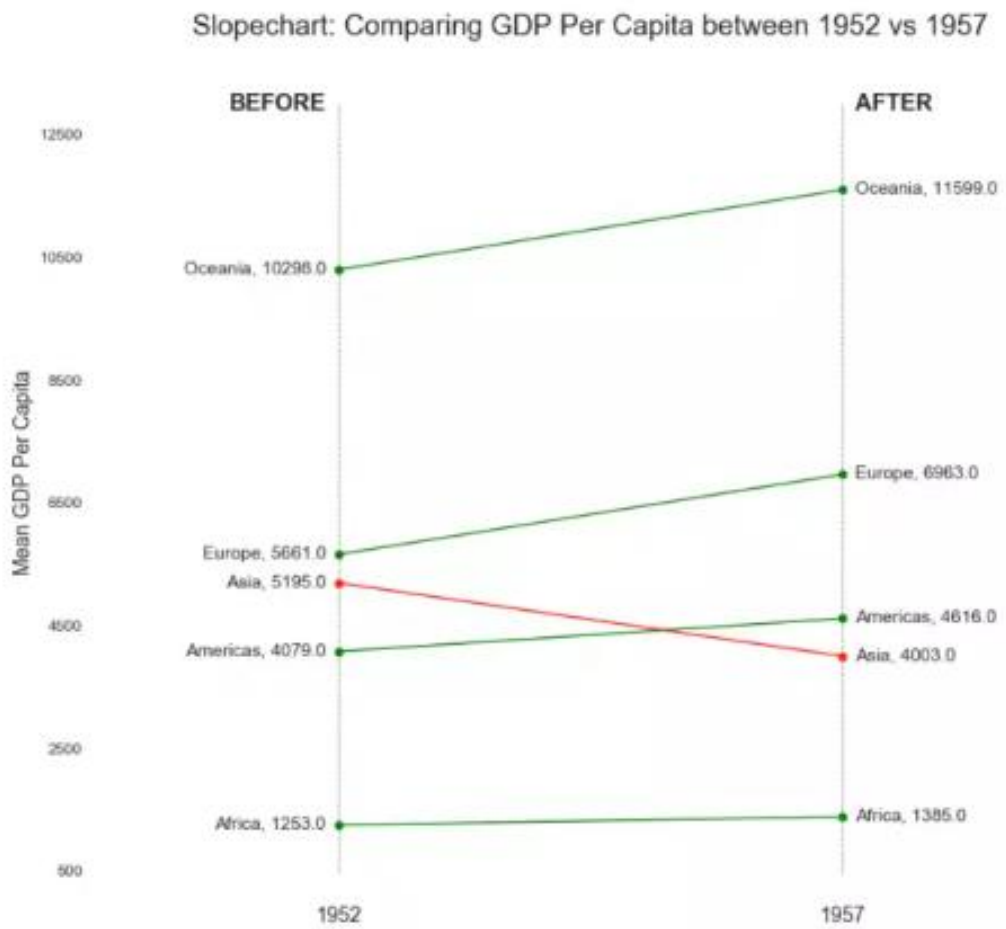

Useful resource

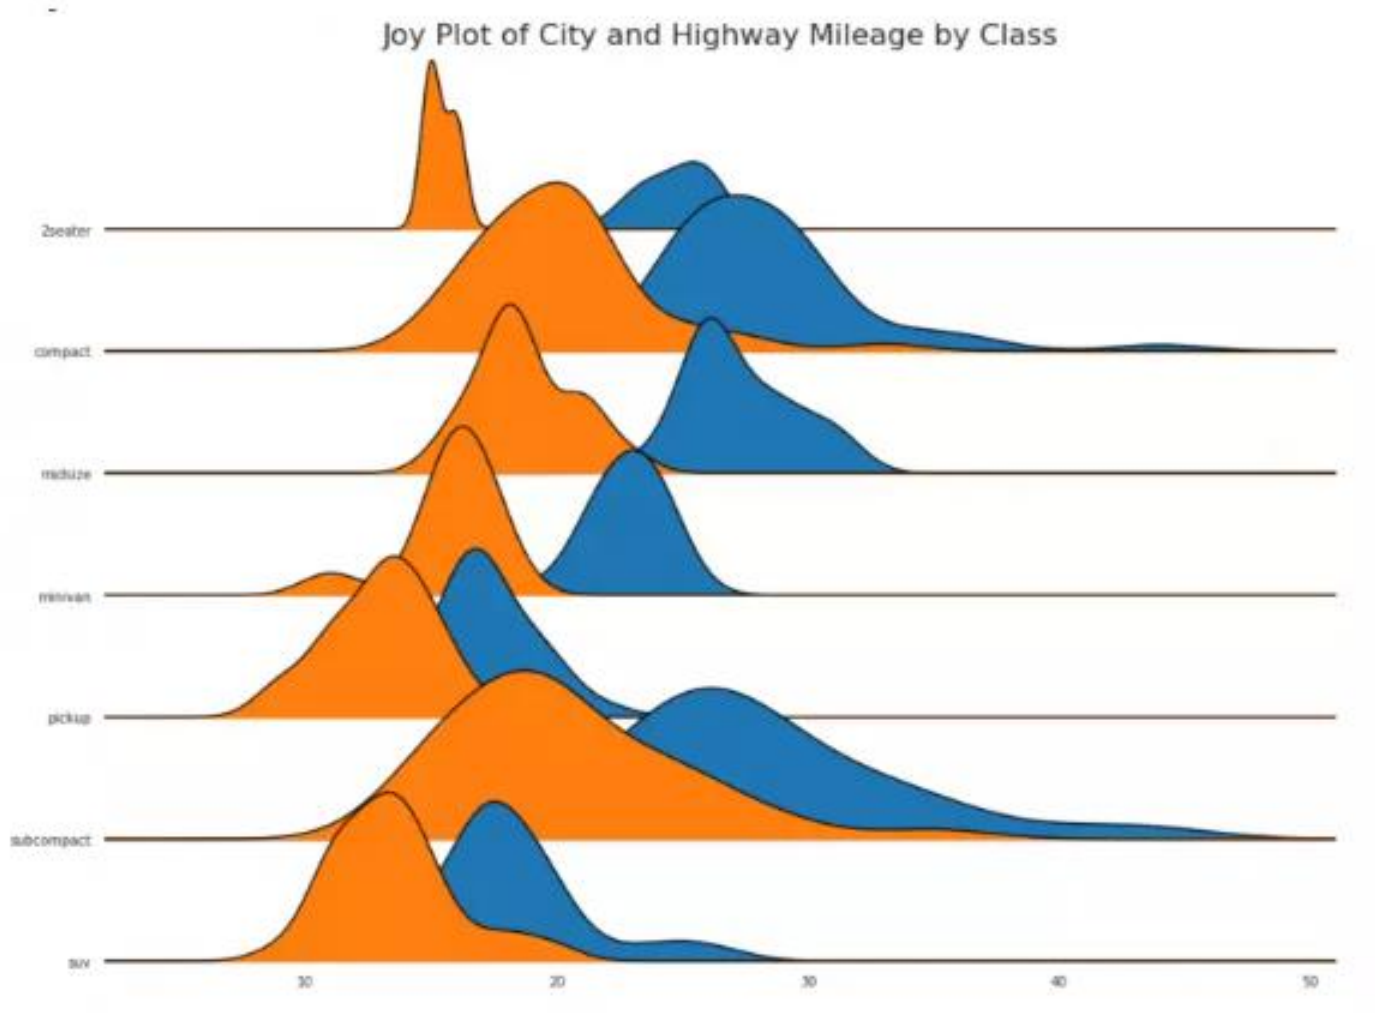

Useful resource

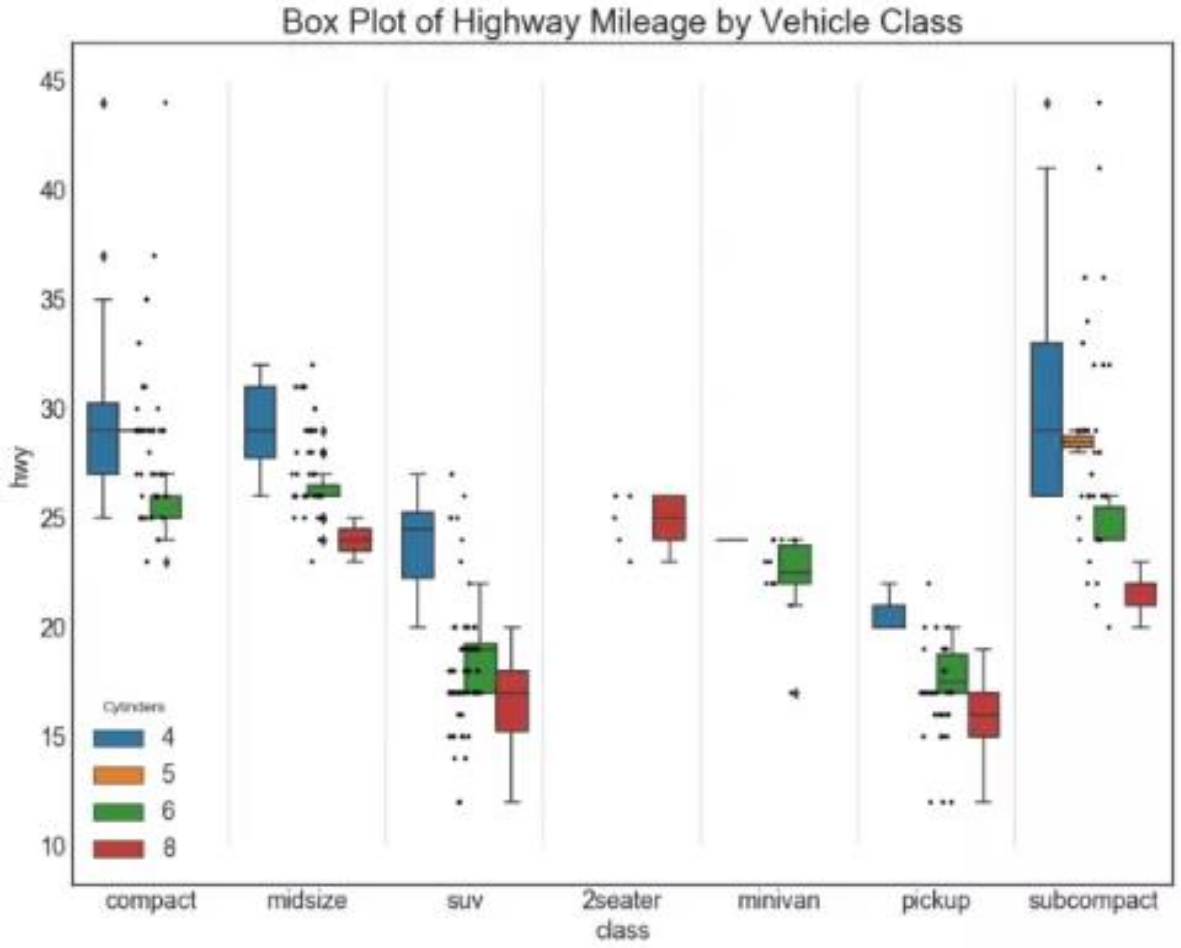

Useful resource

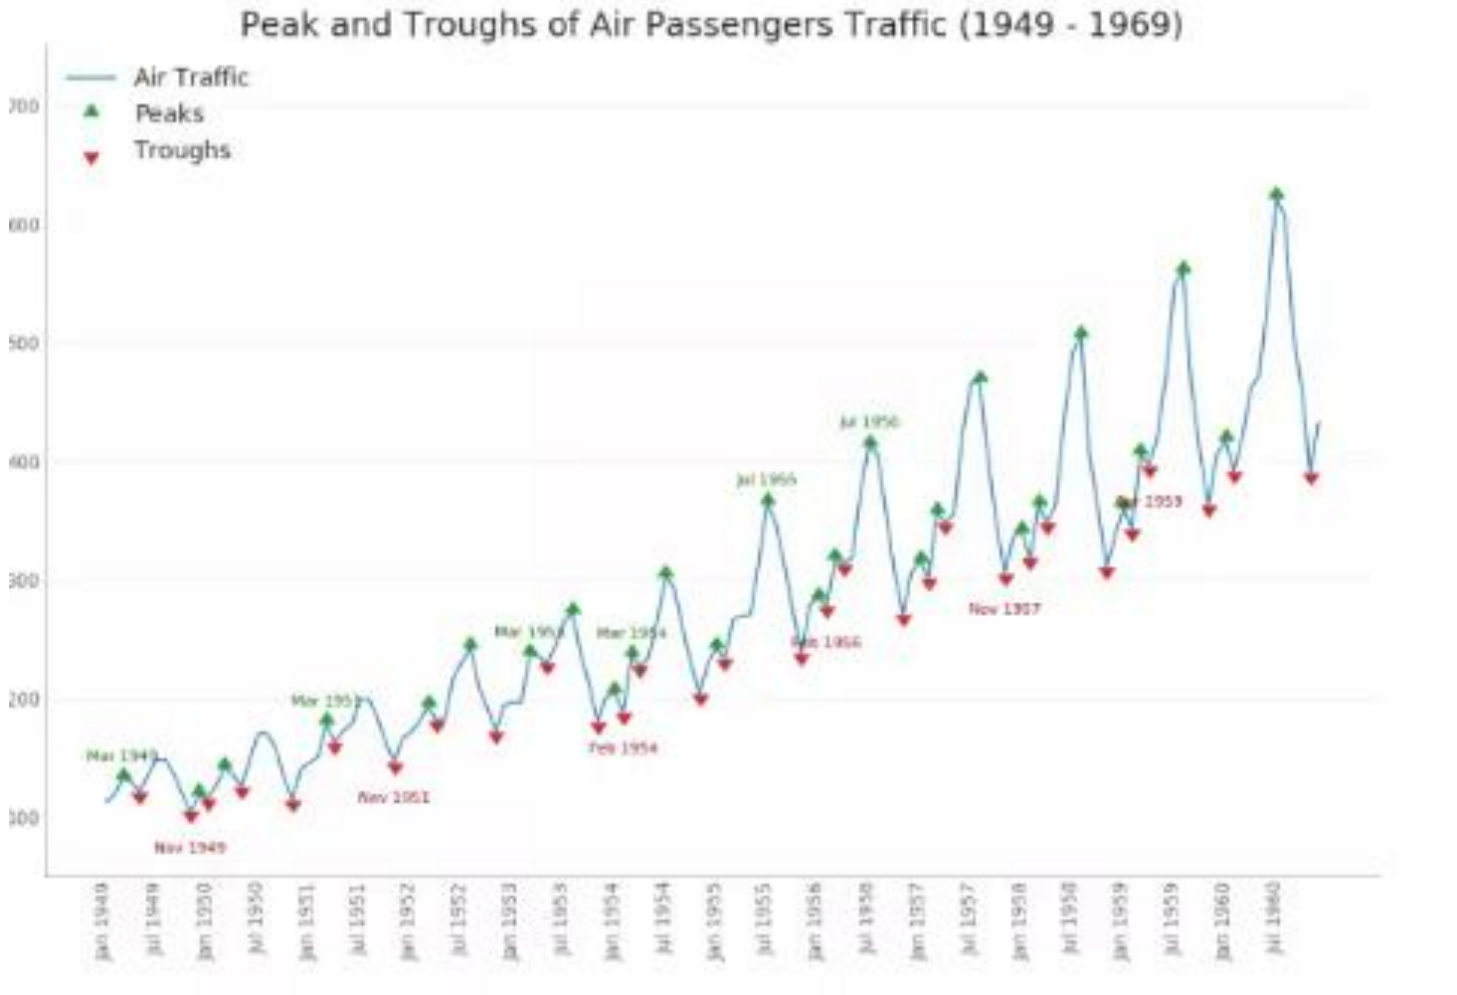

Useful resource

31. Waffle Chart

The `waffle` chart can be created using the `pywaffle` package and is used to show the compositions of groups in a larger population.

Show Code

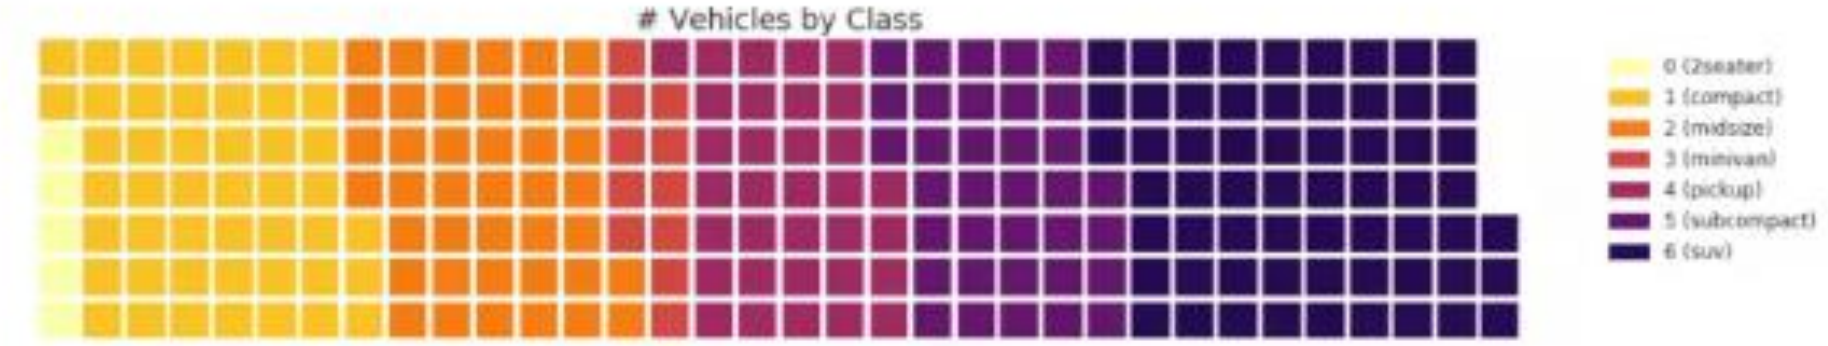

Useful resource

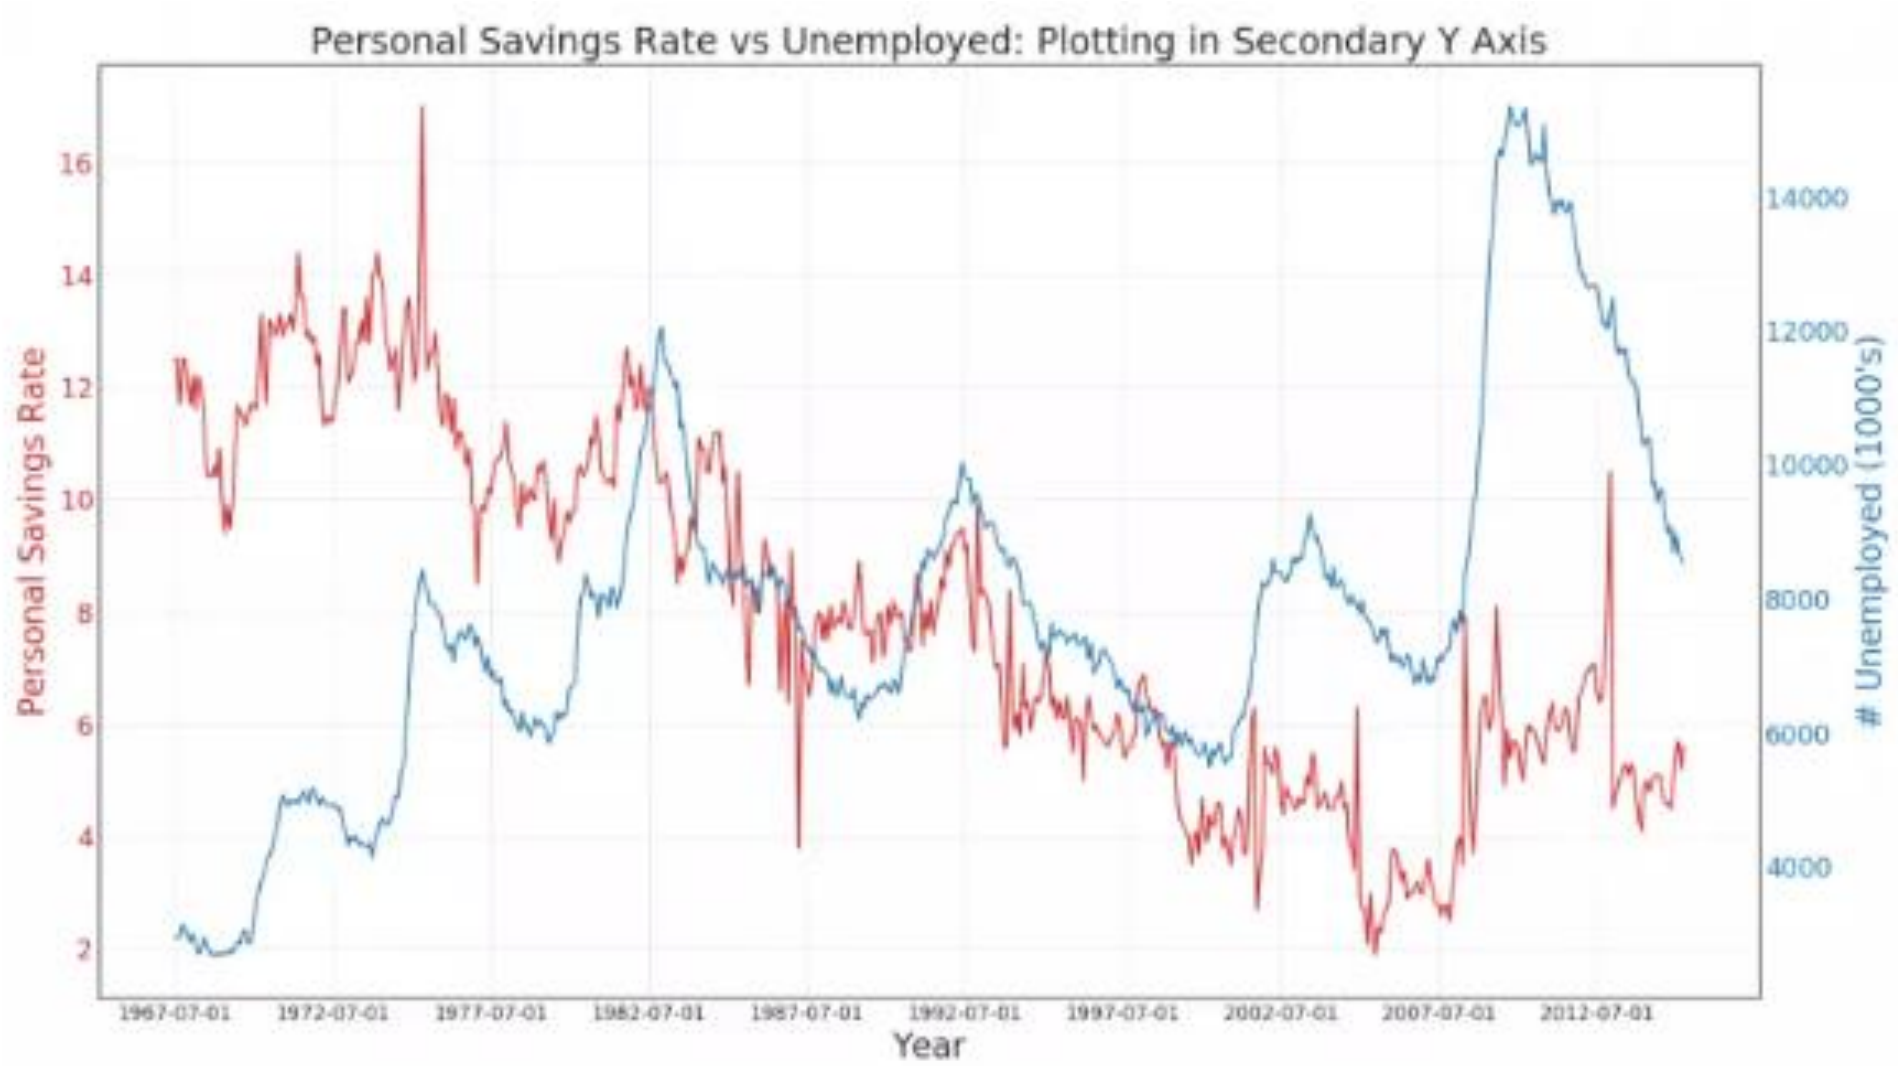

Useful resource

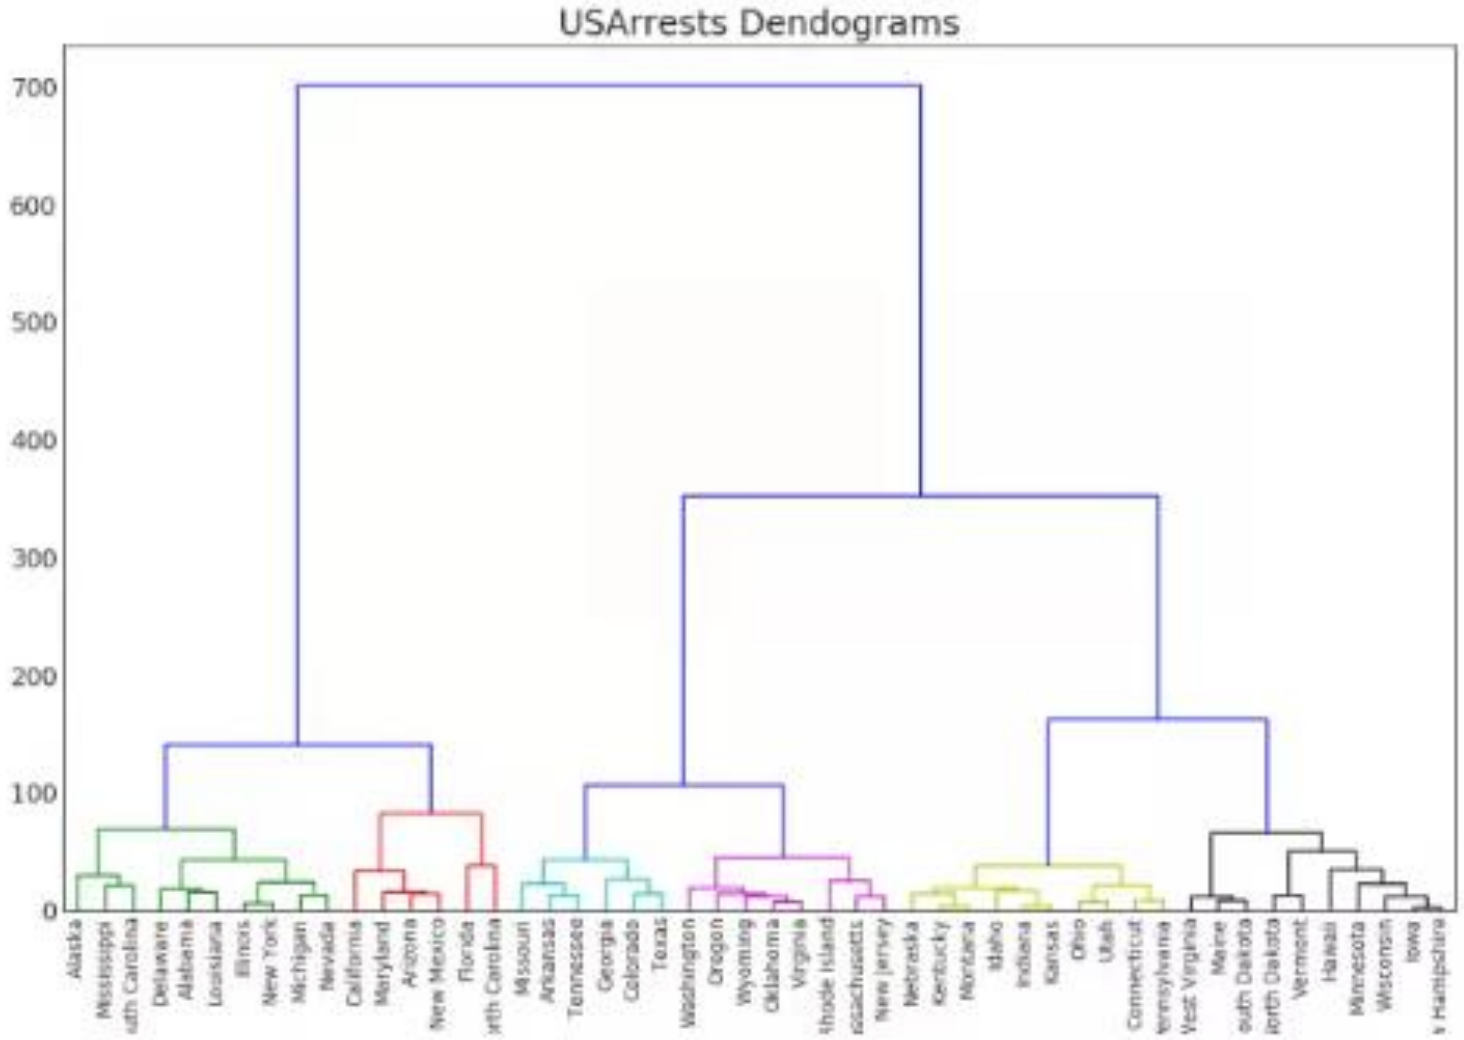

Useful

- https://matplotlib.org/gallery.html

Originally from Genson

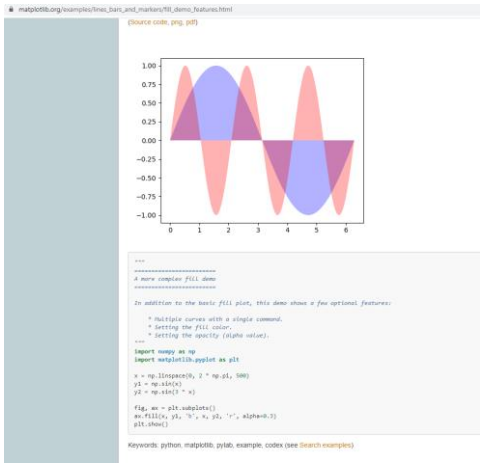

Comparisons

And the winner is simple imputation:

```
MAE from Approach 1 (Drop columns with missing values):
RandomForestRegressor
183550.22137772635

MAE from Approach 2 (Imputation):
RandomForestRegressor
178166.46269899711

MAE from Approach 3 (An Extension to Imputation):
RandomForestRegressor
178927.503183954
```

```
Out[172]: Text(0.5, 1.0, 'Comparision of cleaning options')
```

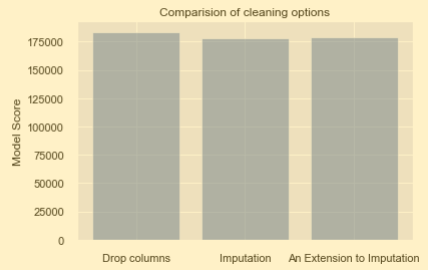


```
In [193]: plt.ylim(170000, 185000)
          plt.bar(options, performance, align='center',  alpha=0.5)
          plt.ylabel('Model Score')

          plt.title('Comparision of cleaning options')
```

```
Out[193]: Text(0.5, 1.0, 'Comparision of cleaning options')
```

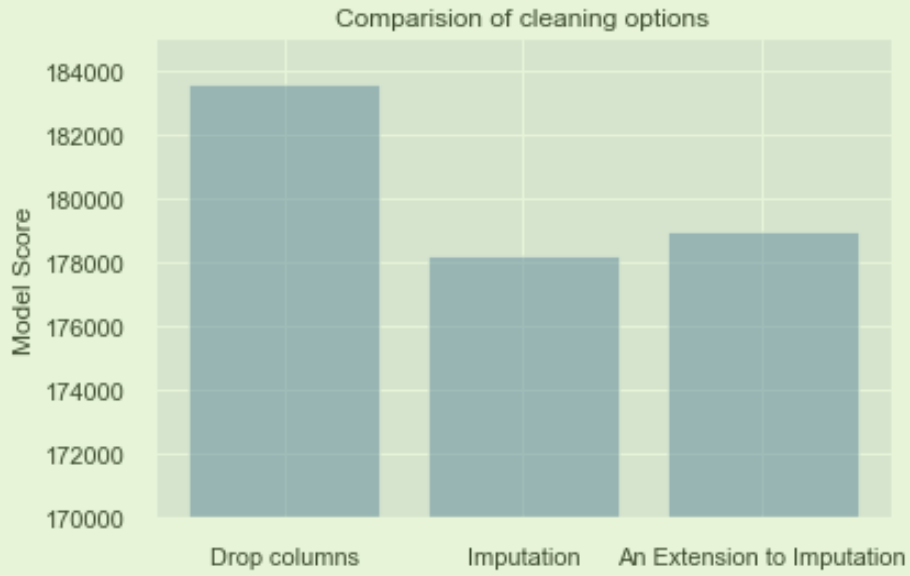

Thanks

Alex Dance

Background

- Maths / statistics degree
- Background in big data, strategy, analytics
- Worked at Optus, Salmat, Reuters, Pathfinder Solutions

Copy of This Presentation and code

https://github.com/alexdance2468/

Plus other data science projects completed

Contact Details

www.linkedin.com/in/alex-dance/

Thanks

Sources:

https://www.kaggle.com/alexisbcook/missing-values - main code for comparisons
https://www.kaggle.com/dansbecker/melbourne-housing-snapshot/home data
https://matplotlib.org/3.1.0/tutorials/text/text_intro.html printing above and below graphs
https://matplotlib.org/gallery.html
https://www.machinelearningplus.com/plots/top-50-matplotlib-visualizations-the-master-plots-python/